

# Temporal\_Difference

April 2, 2018

## 1 Mini Project: Temporal-Difference Methods

In this notebook, you will write your own implementations of many Temporal-Difference (TD) methods.

While we have provided some starter code, you are welcome to erase these hints and write your code from scratch.

### 1.0.1 Part 0: Explore CliffWalkingEnv

Use the code cell below to create an instance of the [CliffWalking](#) environment.

```
In [1]: import gym
        env = gym.make('CliffWalking-v0')
```

The agent moves through a  $4 \times 12$  gridworld, with states numbered as follows:

```
[[ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11],
 [12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23],
 [24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35],
 [36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47]]
```

At the start of any episode, state 36 is the initial state. State 47 is the only terminal state, and the cliff corresponds to states 37 through 46.

The agent has 4 potential actions:

```
UP = 0
RIGHT = 1
DOWN = 2
LEFT = 3
```

Thus,  $\mathcal{S}^+ = \{0, 1, \dots, 47\}$ , and  $\mathcal{A} = \{0, 1, 2, 3\}$ . Verify this by running the code cell below.

```
In [2]: print(env.action_space)
        print(env.observation_space)
```

```
Discrete(4)
Discrete(48)
```

In this mini-project, we will build towards finding the optimal policy for the CliffWalking environment. The optimal state-value function is visualized below. Please take the time now to make sure that you understand *why* this is the optimal state-value function.

```
In [3]: import numpy as np
        from plot_utils import plot_values

        # define the optimal state-value function
        V_opt = np.zeros((4,12))
        V_opt[0:13][0] = -np.arange(3, 15)[::-1]
        V_opt[0:13][1] = -np.arange(3, 15)[::-1] + 1
        V_opt[0:13][2] = -np.arange(3, 15)[::-1] + 2
        V_opt[3][0] = -13

        plot_values(V_opt)

<matplotlib.figure.Figure at 0x7fed29cac400>
```

## 1.0.2 Part 1: TD Prediction: State Values

In this section, you will write your own implementation of TD prediction (for estimating the state-value function).

We will begin by investigating a policy where the agent moves: - RIGHT in states 0 through 10, inclusive,  
- DOWN in states 11, 23, and 35, and - UP in states 12 through 22, inclusive, states 24 through 34, inclusive, and state 36.

The policy is specified and printed below. Note that states where the agent does not choose an action have been marked with -1.

```
In [4]: policy = np.hstack([1*np.ones(11), 2, 0, np.zeros(10), 2, 0, np.zeros(10), 2, 0, -1*np.o
        print("\nPolicy (UP = 0, RIGHT = 1, DOWN = 2, LEFT = 3, N/A = -1):")
        print(policy.reshape(4,12))
```

```
Policy (UP = 0, RIGHT = 1, DOWN = 2, LEFT = 3, N/A = -1):
[[ 1.  1.  1.  1.  1.  1.  1.  1.  1.  1.  1.  2.]
 [ 0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  2.]
 [ 0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  0.  2.]
 [ 0. -1. -1. -1. -1. -1. -1. -1. -1. -1. -1. -1.]]
```

Run the next cell to visualize the state-value function that corresponds to this policy. Make sure that you take the time to understand why this is the corresponding value function!

```
In [5]: V_true = np.zeros((4,12))
        for i in range(3):
            V_true[0:12][i] = -np.arange(3, 15)[::-1] - i
        V_true[1][11] = -2
```

```
V_true[2][11] = -1
V_true[3][0] = -17
```

```
plot_values(V_true)
```



The above figure is what you will try to approximate through the TD prediction algorithm.

Your algorithm for TD prediction has five arguments: - `env`: This is an instance of an OpenAI Gym environment. - `num_episodes`: This is the number of episodes that are generated through agent-environment interaction. - `policy`: This is a 1D numpy array with `policy.shape` equal to the number of states (`env.nS`). `policy[s]` returns the action that the agent chooses when in state `s`. - `alpha`: This is the step-size parameter for the update step. - `gamma`: This is the discount rate. It must be a value between 0 and 1, inclusive (default value: 1).

The algorithm returns as output: - `V`: This is a dictionary where `V[s]` is the estimated value of state `s`.

Please complete the function in the code cell below.

```
In [6]: from collections import defaultdict, deque
import sys

def td_prediction(env, num_episodes, policy, alpha, gamma=1.0):
    # initialize empty dictionaries of floats
    V = defaultdict(float)
    # loop over episodes
    for i_episode in range(1, num_episodes+1):
        # monitor progress
        if i_episode % 100 == 0:
            print("\rEpisode {}/{}".format(i_episode, num_episodes), end="")
            sys.stdout.flush()
        # set initial state
        state = env.reset()

        ## TODO: complete the function
        while True:
            action = policy[state];
```

```

        next_state, reward, done, info = env.step(action);
        V[state] = V[state] + alpha*(reward+gamma*V[next_state] - V[state]);
        if next_state == 47:
            break;
        state = next_state;

    return V

```

Run the code cell below to test your implementation and visualize the estimated state-value function. If the code cell returns **PASSED**, then you have implemented the function correctly! Feel free to change the `num_episodes` and `alpha` parameters that are supplied to the function. However, if you'd like to ensure the accuracy of the unit test, please do not change the value of `gamma` from the default.

```
In [7]: import check_test
```

```
# evaluate the policy and reshape the state-value function
```

```
V_pred = td_prediction(env, 5000, policy, .01)
```

```
# please do not change the code below this line
```


```
V_pred_plot = np.reshape([V_pred[key] if key in V_pred else 0 for key in np.arange(48)],
```

```
check_test.run_check('td_prediction_check', V_pred_plot)
```

```
plot_values(V_pred_plot)
```

Episode 5000/5000

**PASSED**



How close is your estimated state-value function to the true state-value function corresponding to the policy?

You might notice that some of the state values are not estimated by the agent. This is because under this policy, the agent will not visit all of the states. In the TD prediction algorithm, the agent can only estimate the values corresponding to states that are visited.

### 1.0.3 Part 2: TD Control: Sarsa

In this section, you will write your own implementation of the Sarsa control algorithm.

Your algorithm has four arguments: - env: This is an instance of an OpenAI Gym environment.  
- num\_episodes: This is the number of episodes that are generated through agent-environment interaction. - alpha: This is the step-size parameter for the update step. - gamma: This is the discount rate. It must be a value between 0 and 1, inclusive (default value: 1).

The algorithm returns as output: - Q: This is a dictionary (of one-dimensional arrays) where  $Q[s][a]$  is the estimated action value corresponding to state  $s$  and action  $a$ .

Please complete the function in the code cell below.

(Feel free to define additional functions to help you to organize your code.)

```
In [8]: # epsilon-greedy method
def get_probs(Qs, epsilon, nA):
    policy_s = np.ones(nA) * epsilon/nA
    b_action = np.argmax(Qs);
    policy_s[b_action] += 1 - epsilon;
    return policy_s

In [9]: def sarsa(env, num_episodes, alpha, gamma=1.0):
    # initialize action-value function (empty dictionary of arrays)
    Q = defaultdict(lambda: np.zeros(env.nA))
    # initialize performance monitor
    # loop over episodes
    for i_episode in range(1, num_episodes+1):
        # monitor progress
        if i_episode % 100 == 0:
            print("\rEpisode {}/{}".format(i_episode, num_episodes), end="")
            sys.stdout.flush()

        ## TODO: complete the function
        # set initial state
        epsilon = alpha;
        state = env.reset();
        probs = get_probs(Q[state], epsilon, env.nA);
        action = np.random.choice(np.arange(env.nA), p=probs)\
            if state in Q else env.action_space.sample();

        while True:
            next_state, reward, done, info = env.step(action);

            probs = get_probs(Q[next_state], epsilon, env.nA);
            next_action = np.random.choice(np.arange(env.nA), p=probs)\
                if state in Q else env.action_space.sample();
            Q[state][action] += alpha*(reward + gamma*Q[next_state][next_action] - Q[state][action])

            state = next_state;
            action = next_action;

        # whether the agent reaches the terminal state.
```

```

        if done:
            break
    return Q

```

Use the next code cell to visualize the *estimated* optimal policy and the corresponding state-value function.

If the code cell returns **PASSED**, then you have implemented the function correctly! Feel free to change the num\_episodes and alpha parameters that are supplied to the function. However, if you'd like to ensure the accuracy of the unit test, please do not change the value of gamma from the default.

```

In [10]: # obtain the estimated optimal policy and corresponding action-value function
Q_sarsa = sarsa(env, 5000, .01)

# print the estimated optimal policy
policy_sarsa = np.array([np.argmax(Q_sarsa[key]) if key in Q_sarsa else -1 for key in mdp.keys()])
check_test.run_check('td_control_check', policy_sarsa)
print("\nEstimated Optimal Policy (UP = 0, RIGHT = 1, DOWN = 2, LEFT = 3, N/A = -1):")
print(policy_sarsa)

# plot the estimated optimal state-value function
V_sarsa = ([np.max(Q_sarsa[key]) if key in Q_sarsa else 0 for key in mdp.keys()])
plot_values(V_sarsa)

```

Episode 5000/5000


**PASSED**

Estimated Optimal Policy (UP = 0, RIGHT = 1, DOWN = 2, LEFT = 3, N/A = -1):

```

[[ 1  1  2  3  0  1  2  1  2  1  1  2]
 [ 1  0  0  1  0  1  1  1  1  1  2  2]
 [ 1  1  1  1  1  1  1  1  1  1  1  2]
 [ 0 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1  0]]

```



### 1.0.4 Part 3: TD Control: Q-learning

In this section, you will write your own implementation of the Q-learning control algorithm.

Your algorithm has four arguments: - `env`: This is an instance of an OpenAI Gym environment.  
- `num_episodes`: This is the number of episodes that are generated through agent-environment interaction. - `alpha`: This is the step-size parameter for the update step. - `gamma`: This is the discount rate. It must be a value between 0 and 1, inclusive (default value: 1).

The algorithm returns as output: - `Q`: This is a dictionary (of one-dimensional arrays) where `Q[s][a]` is the estimated action value corresponding to state `s` and action `a`.

Please complete the function in the code cell below.

(Feel free to define additional functions to help you to organize your code.)

```
In [11]: def q_learning(env, num_episodes, alpha, gamma=1.0):
    # initialize empty dictionary of arrays
    Q = defaultdict(lambda: np.zeros(env.nA))
    # loop over episodes
    for i_episode in range(1, num_episodes+1):
        # monitor progress
        if i_episode % 100 == 0:
            print("\rEpisode {}/{}".format(i_episode, num_episodes), end="")
            sys.stdout.flush()

        ## TODO: complete the function
        epsilon = alpha;
        state = env.reset();

        while True:
            # take action first
            probs = get_probs(Q[state], epsilon, env.nA);
            action = np.random.choice(np.arange(env.nA), p=probs)\
                if state in Q else env.action_space.sample();
            next_state, reward, done, info = env.step(action);
            best_action = np.argmax(Q[next_state]);

            Q[state][action] += alpha*(reward + gamma*Q[next_state][best_action] - Q[state][action])

            state = next_state;

        # whether the agent reaches the terminal state.
        if done:
            break
    return Q
```

Use the next code cell to visualize the *estimated* optimal policy and the corresponding state-value function.

If the code cell returns **PASSED**, then you have implemented the function correctly! Feel free to change the `num_episodes` and `alpha` parameters that are supplied to the function. However, if you'd like to ensure the accuracy of the unit test, please do not change the value of `gamma` from the default.

```
In [12]: # obtain the estimated optimal policy and corresponding action-value function
Q_sarsamax = q_learning(env, 5000, .01)

# print the estimated optimal policy
policy_sarsamax = np.array([np.argmax(Q_sarsamax[key]) if key in Q_sarsamax else -1 for
check_test.run_check('td_control_check', policy_sarsamax)
print("\nEstimated Optimal Policy (UP = 0, RIGHT = 1, DOWN = 2, LEFT = 3, N/A = -1):")
print(policy_sarsamax)

# plot the estimated optimal state-value function
plot_values([np.max(Q_sarsamax[key]) if key in Q_sarsamax else 0 for key in np.arange(4
```

Episode 5000/5000

**PASSED**

Estimated Optimal Policy (UP = 0, RIGHT = 1, DOWN = 2, LEFT = 3, N/A = -1):

```
[[ 1  0  1  1  1  3  1  0  3  2  2  2]
 [ 1  1  1  1  1  1  1  1  1  2  1  2]
 [ 1  1  1  1  1  1  1  1  1  1  1  2]
 [ 0 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1  0]]
```



### 1.0.5 Part 4: TD Control: Expected Sarsa

In this section, you will write your own implementation of the Expected Sarsa control algorithm.

Your algorithm has four arguments: - env: This is an instance of an OpenAI Gym environment.  
- num\_episodes: This is the number of episodes that are generated through agent-environment interaction. - alpha: This is the step-size parameter for the update step. - gamma: This is the discount rate. It must be a value between 0 and 1, inclusive (default value: 1).

The algorithm returns as output: - Q: This is a dictionary (of one-dimensional arrays) where  $Q[s][a]$  is the estimated action value corresponding to state  $s$  and action  $a$ .

Please complete the function in the code cell below.

(Feel free to define additional functions to help you to organize your code.)



```

In [13]: print(get_probs(Q_sarsa[0], 0.05, env.nA))
          print(Q_sarsa[0])
          print(np.dot(Q_sarsa[0], get_probs(Q_sarsa[0], 0.05, env.nA)))

[ 0.0125  0.9625  0.0125  0.0125]
[-11.16745639 -11.15507427 -11.16038083 -11.16194501]
-11.1553812672

In [14]: def expected_sarsa(env, num_episodes, alpha, gamma=1.0):
          # initialize action-value function (empty dictionary of arrays)
          Q = defaultdict(lambda: np.zeros(env.nA))

          # loop over episodes
          for i_episode in range(1, num_episodes+1):
              # monitor progress
              if i_episode % 100 == 0:
                  print("\rEpisode {}/{}".format(i_episode, num_episodes), end="")

              epsilon = 0.0005;
              # begin an episode
              state = env.reset()
              # get epsilon-greedy action probabilities
              policy_s = get_probs(Q[state], epsilon, env.nA);
              # repeat iterations
              while True:
                  # Choose action At
                  action = np.random.choice(np.arange(env.nA), p=policy_s);
                  # Take action At
                  next_state, reward, done, info = env.step(action);

                  policy_s = get_probs(Q[next_state], epsilon, env.nA);

                  # update Q based on Expected Sarsa
                  Q[state][action] = Q[state][action] + alpha* (reward+ np.dot(Q[next_state],

                  # move to the next state
                  state = next_state;
                  # whether the agent reaches the terminal state.
                  if done:
                      break

          return Q

```

Use the next code cell to visualize the *estimated* optimal policy and the corresponding state-value function.

If the code cell returns **PASSED**, then you have implemented the function correctly! Feel free to change the `num_episodes` and `alpha` parameters that are supplied to the function. However, if you'd like to ensure the accuracy of the unit test, please do not change the value of `gamma` from the default.

```

In [15]: # obtain the estimated optimal policy and corresponding action-value function
Q_expsarsa = expected_sarsa(env, 10000, 1)

# print the estimated optimal policy
policy_expsarsa = np.array([np.argmax(Q_expsarsa[key]) if key in Q_expsarsa else -1 for
check_test.run_check('td_control_check', policy_expsarsa)
print("\nEstimated Optimal Policy (UP = 0, RIGHT = 1, DOWN = 2, LEFT = 3, N/A = -1):")
print(policy_expsarsa)

# plot the estimated optimal state-value function
plot_values([np.max(Q_expsarsa[key]) if key in Q_expsarsa else 0 for key in np.arange(4

```

Episode 10000/10000

**PASSED**

Estimated Optimal Policy (UP = 0, RIGHT = 1, DOWN = 2, LEFT = 3, N/A = -1):

```

[[ 1  0  1  0  1  1  2  2  1  1  1  2]
 [ 2  1  1  1  1  1  1  1  1  1  1  2]
 [ 1  1  1  1  1  1  1  1  1  1  1  2]
 [ 0 -1 -1 -1 -1 -1 -1 -1 -1 -1 -1  0]]

```

